

NN_with_preprocess

December 10, 2015

```
In [186]: import os
import matplotlib.pyplot as plt
%pylab inline
import numpy as np
from lasagne.layers import DenseLayer
from lasagne.layers import InputLayer
from lasagne.layers import DropoutLayer
from lasagne.layers import Conv2DLayer
from lasagne.layers import MaxPool2DLayer
from lasagne.nonlinearities import softmax
from lasagne.updates import adam
from lasagne.layers import get_all_params
from nolearn.lasagne import NeuralNet
from nolearn.lasagne import TrainSplit
from nolearn.lasagne import objective
```

Populating the interactive namespace from numpy and matplotlib

WARNING: pylab import has clobbered these variables: ['axes', 'f']
'%matplotlib' prevents importing * from pylab and numpy

```
In [187]: import scipy.io
'''
train = scipy.io.loadmat('labeled_images.mat')
print "Shape of tr_images is: ", train["tr_images"].shape
(x_size, y_size, n_images) = train["tr_images"].shape
X = np.reshape(np.swapaxes(train["tr_images"], 0, 2), (n_images, 1, x_size, y_size))
y = train["tr_labels"].ravel()-1
print X.shape
print y.shape
X = np.array(X).astype(np.float32)
y = np.array(y).astype(np.int32)
# Normalization
X -= X.mean()
X /= X.std()
print X[0].shape
print y.shape
'''

train = scipy.io.loadmat('filtered_testing.mat')
train_original = scipy.io.loadmat('labeled_images.mat')

print "Shape of tr_images is: ", train["tr_images"].shape
# (x, y, n_images) = train["tr_images"].shape
```

```

(n_images, dim) = train["tr_images"].shape
y = train_original["tr_labels"].ravel()-1

# train_img = np.reshape(np.swapaxes(train["tr_images"], 0, 2), (n_images, x * y))
X = np.reshape(train['tr_images'], (n_images, 1, dim, 1))
y = np.array(y).astype(np.int32)
X = np.array(X).astype(np.float32)
#X -= X.mean()
#X /= X.std()

#plt.imshow(np.swapaxes(np.reshape(train_img[0], (y, x)), 0, 1), cmap=pylab.gray())
#plt.show()

#plt.imshow(np.swapaxes(np.reshape(X[0], (y_size, x_size)), 0, 1), cmap=pylab.gray())
#plt.show()

```

Shape of tr_images is: (2925, 2560)

```

In [188]: # Show labels of the dataset
figs, axes = plt.subplots(4, 4, figsize=(6, 6))
for i in range(4):
    for j in range(4):
        axes[i, j].imshow(-X[i + 4 * j].reshape(32, 32), cmap='gray', interpolation='none')
        axes[i, j].set_xticks([])
        axes[i, j].set_yticks([])
        axes[i, j].set_title("Label: {}".format(y[i + 4 * j]))
        axes[i, j].axis('off')

```

ValueError


Traceback (most recent call last)

```

<ipython-input-188-97195b50896d> in <module>()
      3 for i in range(4):
      4     for j in range(4):
----> 5         axes[i, j].imshow(-X[i + 4 * j].reshape(32, 32), cmap='gray', interpolation='none')
      6         axes[i, j].set_xticks([])
      7         axes[i, j].set_yticks([])

```

ValueError: total size of new array must be unchanged



```
In [189]: layers0 = [
    # layer dealing with the input data
    (InputLayer, {'shape': (None, X.shape[1], X.shape[2], X.shape[3])}),

    # first stage of our convolutional layers
    (Conv2DLayer, {'num_filters': 96, 'filter_size': 5}),
    (Conv2DLayer, {'num_filters': 96, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 96, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 96, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 96, 'filter_size': 3}),
    (MaxPool2DLayer, {'pool_size': 2}),

    # second stage of our convolutional layers
    (Conv2DLayer, {'num_filters': 128, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 128, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 128, 'filter_size': 3}),
    (MaxPool2DLayer, {'pool_size': 2}),

    # two dense layers with dropout
    (DenseLayer, {'num_units': 64}),
```

```

(DropoutLayer, {}),
(DenseLayer, {'num_units': 64}),

# the output layer
(DenseLayer, {'num_units': 7, 'nonlinearity': softmax}),
]

layers1 = [
    # layer dealing with the input data
    (InputLayer, {'shape': (None, X.shape[1], X.shape[2], X.shape[3])}),

    # first stage of our convolutional layers
    (Conv2DLayer, {'num_filters': 48, 'filter_size': 5}),
    (Conv2DLayer, {'num_filters': 48, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 48, 'filter_size': 3}),
    (MaxPool2DLayer, {'pool_size': 2}),

    # second stage of our convolutional layers
    (Conv2DLayer, {'num_filters': 64, 'filter_size': 5}),
    (Conv2DLayer, {'num_filters': 64, 'filter_size': 3}),
    (MaxPool2DLayer, {'pool_size': 2}),

    # two dense layers with dropout
    (DenseLayer, {'num_units': 32}),
    (DropoutLayer, {}),
    (DenseLayer, {'num_units': 32}),

    # the output layer
    (DenseLayer, {'num_units': 7, 'nonlinearity': softmax}),
]

layers2 = [
    (InputLayer, {'shape': (None, X.shape[1], X.shape[2], X.shape[3])}),

    (Conv2DLayer, {'num_filters': 32, 'filter_size': (3, 3)}),
    (MaxPool2DLayer, {'pool_size': (2, 2)}),

    (Conv2DLayer, {'num_filters': 64, 'filter_size': (3, 3)}),
    (Conv2DLayer, {'num_filters': 64, 'filter_size': (3, 3)}),
    (MaxPool2DLayer, {'pool_size': (2, 2)}),

    (Conv2DLayer, {'num_filters': 96, 'filter_size': (3, 3)}),
    (MaxPool2DLayer, {'pool_size': (2, 2)}),

    (DenseLayer, {'num_units': 64}),
    (DropoutLayer, {}),
    (DenseLayer, {'num_units': 64}),

    (DenseLayer, {'num_units': 7, 'nonlinearity': softmax}),
]

layers3 = [
    (InputLayer, {'shape': (None, X.shape[1], X.shape[2], X.shape[3])}),
    (Conv2DLayer, {'num_filters': 96, 'filter_size': (5, 5)}),

```

```

(MaxPool2DLayer, {'pool_size': (2, 2)}),
(Conv2DLayer, {'num_filters': 64, 'filter_size': (5, 5)}),
(MaxPool2DLayer, {'pool_size': (2, 2)}),
(DenseLayer, {'num_units': 64}),
(DenseLayer, {'num_units': 7, 'nonlinearity': softmax}),
]

layers4 = [
    (InputLayer, {'shape': (None, X.shape[1], X.shape[2], X.shape[3])}),
    #(Conv2DLayer, {'num_filters': 96, 'filter_size': (5, 5)}),
    #(MaxPool2DLayer, {'pool_size': (2, 2)}),
    (DenseLayer, {'num_units': 90}),
    (DenseLayer, {'num_units': 7, 'nonlinearity': softmax}),
]

In [190]: def regularization_objective(layers, lambda1=0., lambda2=0., *args, **kwargs):
    # default loss
    losses = objective(layers, *args, **kwargs)
    # get the layers' weights, but only those that should be regularized
    # (i.e. not the biases)
    weights = get_all_params(layers[-1], regularizable=True)
    # sum of absolute weights for L1
    sum_abs_weights = sum([abs(w).sum() for w in weights])
    # sum of squared weights for L2
    sum_squared_weights = sum([(w ** 2).sum() for w in weights])
    # add weights to regular loss
    losses += lambda1 * sum_abs_weights + lambda2 * sum_squared_weights
    return losses

In [191]: net0 = NeuralNet(
    layers=layers0,
    max_epochs=20,

    update=adam,
    update_learning_rate=0.0002,

    objective=regularization_objective,
    objective_lambda2=0.0025,

    train_split=TrainSplit(eval_size=0.25),
    verbose=4,
)
net1 = NeuralNet(
    layers=layers4,
    max_epochs=100,
    update=adam,
    update_learning_rate=0.0001,
    objective=regularization_objective,
    objective_lambda2=0.002,
    train_split=TrainSplit(eval_size=0.005),
    verbose=3,
)

In [192]: net1.fit(X, y)

```

Neural Network with 231127 learnable parameters

Layer information

#	name	size
0	input0	1x2560x1
1	dense1	90
2	dense2	7

epoch	train loss	valid loss	train/val	valid acc	dur
1	2.54987	2.31965	1.09925	0.33333	0.29s
2	1.98719	1.98705	1.00007	0.38889	0.32s
3	1.68112	1.84959	0.90892	0.50000	0.30s
4	1.49158	1.76722	0.84402	0.50000	0.26s
5	1.37405	1.71639	0.80055	0.55556	0.26s
6	1.29418	1.65386	0.78252	0.55556	0.27s
7	1.22849	1.61063	0.76273	0.55556	0.27s
8	1.17063	1.58512	0.73851	0.55556	0.27s
9	1.11278	1.55161	0.71718	0.61111	0.25s
10	1.06248	1.50143	0.70765	0.66667	0.25s
11	1.02617	1.46533	0.70030	0.66667	0.26s
12	0.99616	1.45201	0.68606	0.66667	0.28s
13	0.96675	1.45138	0.66609	0.66667	0.26s
14	0.93497	1.45026	0.64469	0.66667	0.28s
15	0.90458	1.44376	0.62655	0.66667	0.26s
16	0.87694	1.42061	0.61729	0.66667	0.26s
17	0.85375	1.39171	0.61346	0.66667	0.26s
18	0.83461	1.36140	0.61305	0.66667	0.26s
19	0.81777	1.34873	0.60633	0.72222	0.26s
20	0.80239	1.34948	0.59459	0.72222	0.26s
21	0.78649	1.34350	0.58541	0.72222	0.26s
22	0.77115	1.35035	0.57107	0.72222	0.29s
23	0.75470	1.35405	0.55737	0.72222	0.33s
24	0.73940	1.35321	0.54640	0.66667	0.29s
25	0.72426	1.34923	0.53679	0.72222	0.27s
26	0.71023	1.34199	0.52924	0.72222	0.26s
27	0.69861	1.32925	0.52556	0.72222	0.26s
28	0.68784	1.31917	0.52142	0.77778	0.26s
29	0.67754	1.30962	0.51736	0.77778	0.28s
30	0.66778	1.30994	0.50978	0.77778	0.26s
31	0.65762	1.31417	0.50041	0.72222	0.27s
32	0.64770	1.32188	0.48999	0.72222	0.26s
33	0.63855	1.32304	0.48264	0.77778	0.26s
34	0.62958	1.32117	0.47653	0.77778	0.27s
35	0.62078	1.31987	0.47033	0.72222	0.28s
36	0.61189	1.31582	0.46502	0.77778	0.32s
37	0.60358	1.31718	0.45824	0.77778	0.30s
38	0.59537	1.31575	0.45249	0.77778	0.27s
39	0.58779	1.30946	0.44888	0.77778	0.31s
40	0.58097	1.29817	0.44753	0.77778	0.29s
41	0.57407	1.29932	0.44182	0.77778	0.28s
42	0.56818	1.29923	0.43732	0.77778	0.27s

43	0.56227	1.30237	0.43173	0.77778	0.28s
44	0.55659	1.29413	0.43009	0.77778	0.27s
45	0.55061	1.30487	0.42196	0.77778	0.36s
46	0.54444	1.31491	0.41405	0.77778	0.47s
47	0.53802	1.32319	0.40661	0.77778	0.38s
48	0.53207	1.32608	0.40123	0.77778	0.29s
49	0.52585	1.31750	0.39913	0.77778	0.29s
50	0.52097	1.31233	0.39698	0.72222	0.31s
51	0.51525	1.31774	0.39101	0.72222	0.27s
52	0.51036	1.31170	0.38908	0.72222	0.26s
53	0.50528	1.30091	0.38841	0.72222	0.26s
54	0.50131	1.29803	0.38621	0.72222	0.27s
55	0.49707	1.30047	0.38222	0.72222	0.27s
56	0.49305	1.30491	0.37784	0.72222	0.26s
57	0.48926	1.30371	0.37528	0.72222	0.27s
58	0.48573	1.30207	0.37305	0.72222	0.28s
59	0.48136	1.31143	0.36705	0.72222	0.26s
60	0.47721	1.32565	0.35998	0.72222	0.26s
61	0.47286	1.33330	0.35465	0.72222	0.27s
62	0.46869	1.33236	0.35178	0.72222	0.29s
63	0.46460	1.33348	0.34841	0.72222	0.25s
64	0.46035	1.33734	0.34422	0.72222	0.25s
65	0.45642	1.33358	0.34225	0.72222	0.27s
66	0.45271	1.32799	0.34090	0.72222	0.26s
67	0.44915	1.32018	0.34021	0.72222	0.26s
68	0.44574	1.31147	0.33988	0.72222	0.26s
69	0.44266	1.31206	0.33737	0.72222	0.28s
70	0.44037	1.30611	0.33716	0.72222	0.26s
71	0.43811	1.30204	0.33648	0.72222	0.26s
72	0.43563	1.29472	0.33646	0.72222	0.27s
73	0.43321	1.30359	0.33232	0.72222	0.37s
74	0.43065	1.32220	0.32571	0.72222	0.44s
75	0.42791	1.34024	0.31928	0.72222	0.27s
76	0.42459	1.34718	0.31517	0.72222	0.26s
77	0.42128	1.35974	0.30982	0.72222	0.29s
78	0.41753	1.36796	0.30522	0.72222	0.26s
79	0.41407	1.36723	0.30285	0.72222	0.26s
80	0.41120	1.36049	0.30225	0.72222	0.26s
81	0.40845	1.34545	0.30358	0.72222	0.26s
82	0.40583	1.33576	0.30382	0.72222	0.28s
83	0.40304	1.30944	0.30780	0.72222	0.26s
84	0.40078	1.29022	0.31063	0.72222	0.25s
85	0.39902	1.27397	0.31321	0.72222	0.29s
86	0.39821	1.27082	0.31335	0.72222	0.27s
87	0.39769	1.28061	0.31055	0.72222	0.27s
88	0.39713	1.29406	0.30688	0.72222	0.27s
89	0.39624	1.31851	0.30052	0.72222	0.43s
90	0.39426	1.34113	0.29397	0.72222	0.42s
91	0.39128	1.37099	0.28540	0.72222	0.33s
92	0.38738	1.38747	0.27920	0.72222	0.43s
93	0.38362	1.40467	0.27311	0.72222	0.28s
94	0.38063	1.41950	0.26815	0.66667	0.26s
95	0.37927	1.41203	0.26860	0.72222	0.27s
96	0.37842	1.37411	0.27539	0.72222	0.28s

97	0.37614	1.32176	0.28458	0.72222	0.29s
98	0.37249	1.27419	0.29234	0.72222	0.44s
99	0.36935	1.24197	0.29739	0.72222	0.32s
100	0.36916	1.23651	0.29855	0.72222	0.26s

```
Out[192]: NeuralNet(X_tensor_type=None,  
    batch_iterator_test=<nolearn.lasagne.base.BatchIterator object at 0x1077bf150>,  
    batch_iterator_train=<nolearn.lasagne.base.BatchIterator object at 0x1077bf0d0>,  
    custom_score=None,  
    layers=[(<class 'lasagne.layers.input.InputLayer'>, {'shape': (None, 1, 2560, 1)}), (<cl  
    loss=None, max_epochs=100, more_params={},  
    objective=<function regularization_objective at 0x1301b3d70>,  
    objective_lambda2=0.002,  
    objective_loss_function=<function categorical_crossentropy at 0x1074af758>,  
    on_batch_finished=[],  
    on_epoch_finished=[<nolearn.lasagne.handlers.PrintLog instance at 0x1303aa3f8>],  
    on_training_finished=[],  
    on_training_started=[<nolearn.lasagne.handlers.PrintLayerInfo instance at 0x12f543710>],  
    regression=False,  
    train_split=<nolearn.lasagne.base.TrainSplit object at 0x1306b4e50>,  
    update=<function adam at 0x1074b7b18>, update_learning_rate=0.0001,  
    use_label_encoder=False, verbose=3,  
    y_tensor_type=TensorType(int32, vector))
```

```
In [193]: net1.save_params_to ('NN_model_with_preprocess')
```

```
In [194]: def classify_pub_test(classifier):
    '''
    test = scipy.io.loadmat('public_test_images.mat')
    print test
    print test["public_test_images"].shape
    (x, y, n_images) = test["public_test_images"].shape
    test_img = np.reshape(np.swapaxes(test["public_test_images"], 0, 2), (n_images, 1, x, y))

    test_img = np.array(test_img).astype(np.float32)
    test_img -= test_img.mean()
    test_img /= test_img.std()
    '''

    test = scipy.io.loadmat('public_test_filtered_no_normalization.mat')
    pub_test = scipy.io.loadmat('./public_test_filtered_no_normalization.mat')
    hid_test = scipy.io.loadmat('./hidden_test_images_filtered.mat')
    (n_images, dim) = pub_test["public_test_images"].shape
    test_img = np.reshape(pub_test['public_test_images'], (n_images, 1, dim, 1))
    test_img = np.array(test_img).astype(np.float32)
    #test_img -= test_img.mean()
    #test_img /= test_img.std()
    pub_res = list(classifier.predict(test_img)+1)

    (n_images, dim) = hid_test["hidden_img"].shape
    test_img = np.reshape(hid_test["hidden_img"], (n_images, 1, dim, 1))
    test_img = np.array(test_img).astype(np.float32)
    #test_img -= test_img.mean()
    #test_img /= test_img.std()
    hid_res = list(classifier.predict(test_img)+1)
    return pub_res+hid_res
```



```

In [195]: classify_result = classify_pub_test(net1)
          cls_res_list = list(classify_result)
          print cls_res_list
          with open('submit_nn_sing_layer_90_units_100iter_non_normalized.csv', 'w') as f:
              f.write('Id,Prediction\n')
              index = 1
              for pred in cls_res_list:
                  f.write('%d,%d\n'%(index, pred))
                  index += 1
              while index<=1253:
                  f.write('%d,0\n'%(index))
                  index+=1

```

[7, 5, 6, 7, 7, 7, 7, 7, 7, 7, 7, 7, 6, 7, 7, 7, 4, 7, 4, 4, 7, 4, 7, 7, 7, 7, 4, 7, 4, 7, 4, 4, 4, 2, 5]

```

In [ ]: from nolearn.lasagne.visualize import plot_loss
          from nolearn.lasagne.visualize import plot_conv_weights
          from nolearn.lasagne.visualize import plot_conv_activity
          from nolearn.lasagne.visualize import plot_occlusion
          plot_loss(net1)
          plot_conv_activity(net1.layers_[1], X[0:1])

```

In []:

In []:

In []:

In []: