

NN_with_preprocess

December 11, 2015

```
In [221]: import os
import matplotlib.pyplot as plt
%pylab inline
import numpy as np
from lasagne.layers import DenseLayer
from lasagne.layers import InputLayer
from lasagne.layers import DropoutLayer
from lasagne.layers import Conv2DLayer
from lasagne.layers import MaxPool2DLayer
from lasagne.nonlinearities import softmax
from lasagne.updates import adam
from lasagne.layers import get_all_params
from nolearn.lasagne import NeuralNet
from nolearn.lasagne import TrainSplit
from nolearn.lasagne import objective
```

Populating the interactive namespace from numpy and matplotlib

```
In [244]: import scipy.io
```

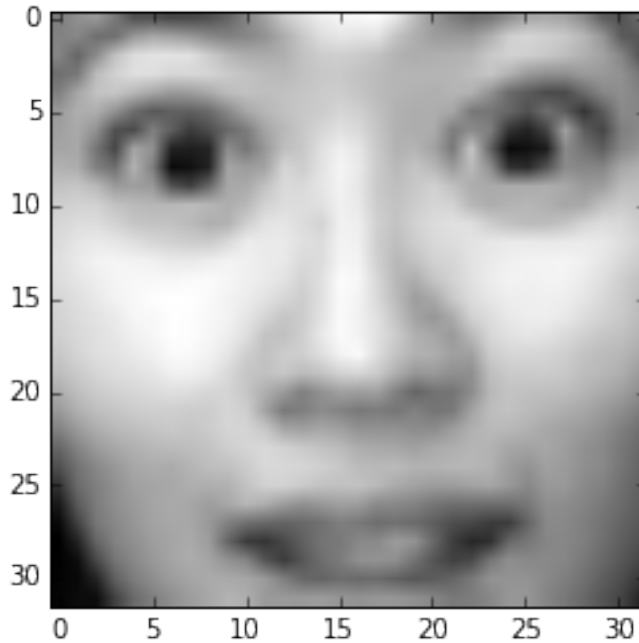
```
train = scipy.io.loadmat('filtered_testing.mat')
train_original = scipy.io.loadmat('labeled_images.mat')
```

```
print "Shape of tr_images is: ", train["tr_images"].shape
# (x, y, n_images) = train["tr_images"].shape
(n_images, dim) = train["tr_images"].shape
y = train_original["tr_labels"].ravel()-1
```

```
# train_img = np.reshape(np.swapaxes(train["tr_images"], 0, 2), (n_images, x * y))
X = np.reshape(train['tr_images'], (n_images, 1, dim, 1))
y = np.array(y).astype(np.int32)
X = np.array(X).astype(np.float32)
# Normalization
#X -= X.mean()
#X /= X.std()
```

```
plt.imshow(np.swapaxes((np.swapaxes(train_original["tr_images"], 0, 2)), 1,2)[0], cmap=pylab.cm.gray)
plt.show()
```

Shape of tr_images is: (2925, 2560)



```
In [245]: layers0 = [
    # layer dealing with the input data
    (InputLayer, {'shape': (None, X.shape[1], X.shape[2], X.shape[3])}),

    # first stage of our convolutional layers
    (Conv2DLayer, {'num_filters': 96, 'filter_size': 5}),
    (Conv2DLayer, {'num_filters': 96, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 96, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 96, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 96, 'filter_size': 3}),
    (MaxPool2DLayer, {'pool_size': 2}),

    # second stage of our convolutional layers
    (Conv2DLayer, {'num_filters': 128, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 128, 'filter_size': 3}),
    (Conv2DLayer, {'num_filters': 128, 'filter_size': 3}),
    (MaxPool2DLayer, {'pool_size': 2}),

    # two dense layers with dropout
    (DenseLayer, {'num_units': 64}),
    (DropoutLayer, {}),
    (DenseLayer, {'num_units': 64}),

    # the output layer
    (DenseLayer, {'num_units': 7, 'nonlinearity': softmax}),
]

layers1 = [
    # layer dealing with the input data
```

```

        (InputLayer, {'shape': (None, X.shape[1], X.shape[2], X.shape[3])}),

        # first stage of our convolutional layers
        (Conv2DLayer, {'num_filters': 48, 'filter_size': 5}),
        (Conv2DLayer, {'num_filters': 48, 'filter_size': 3}),
        (Conv2DLayer, {'num_filters': 48, 'filter_size': 3}),
        (MaxPool2DLayer, {'pool_size': 2}),

        # second stage of our convolutional layers
        (Conv2DLayer, {'num_filters': 64, 'filter_size': 5}),
        (Conv2DLayer, {'num_filters': 64, 'filter_size': 3}),
        (MaxPool2DLayer, {'pool_size': 2}),

        # two dense layers with dropout
        (DenseLayer, {'num_units': 32}),
        (DropoutLayer, {}),
        (DenseLayer, {'num_units': 32}),

        # the output layer
        (DenseLayer, {'num_units': 7, 'nonlinearity': softmax}),
    ]

layers2 = [
    (InputLayer, {'shape': (None, X.shape[1], X.shape[2], X.shape[3])}),

    (Conv2DLayer, {'num_filters': 32, 'filter_size': (3, 3)}),
    (MaxPool2DLayer, {'pool_size': (2, 2)}),

    (Conv2DLayer, {'num_filters': 64, 'filter_size': (3, 3)}),
    (Conv2DLayer, {'num_filters': 64, 'filter_size': (3, 3)}),
    (MaxPool2DLayer, {'pool_size': (2, 2)}),

    (Conv2DLayer, {'num_filters': 96, 'filter_size': (3, 3)}),
    (MaxPool2DLayer, {'pool_size': (2, 2)}),

    (DenseLayer, {'num_units': 64}),
    (DropoutLayer, {}),
    (DenseLayer, {'num_units': 64}),

    (DenseLayer, {'num_units': 7, 'nonlinearity': softmax}),
]

layers3 = [
    (InputLayer, {'shape': (None, X.shape[1], X.shape[2], X.shape[3])}),
    (Conv2DLayer, {'num_filters': 96, 'filter_size': (5, 5)}),
    (MaxPool2DLayer, {'pool_size': (2, 2)}),
    (Conv2DLayer, {'num_filters': 64, 'filter_size': (5, 5)}),
    (MaxPool2DLayer, {'pool_size': (2, 2)}),
    (DenseLayer, {'num_units': 64}),
    (DenseLayer, {'num_units': 7, 'nonlinearity': softmax}),
]

layers4 = [
    (InputLayer, {'shape': (None, X.shape[1], X.shape[2], X.shape[3])}),

```

```

        #(Conv2DLayer, {'num_filters': 96, 'filter_size': (5, 5)}),
        #(MaxPool2DLayer, {'pool_size': (2, 2)}),
        (DenseLayer, {'num_units': 90}),
        (DenseLayer, {'num_units': 7, 'nonlinearity': softmax}),
    ]

```

```

In [246]: def regularization_objective(layers, lambda1=0., lambda2=0., *args, **kwargs):
    # default loss
    losses = objective(layers, *args, **kwargs)
    # get the layers' weights, but only those that should be regularized
    # (i.e. not the biases)
    weights = get_all_params(layers[-1], regularizable=True)
    # sum of absolute weights for L1
    sum_abs_weights = sum([abs(w).sum() for w in weights])
    # sum of squared weights for L2
    sum_squared_weights = sum([(w ** 2).sum() for w in weights])
    # add weights to regular loss
    losses += lambda1 * sum_abs_weights + lambda2 * sum_squared_weights
    return losses

```

```

In [253]: net0 = NeuralNet(
    layers=layers0,
    max_epochs=20,

    update=adam,
    update_learning_rate=0.0002,

    objective=regularization_objective,
    objective_lambda2=0.0025,

    train_split=TrainSplit(eval_size=0.25),
    verbose=4,
)
net1 = NeuralNet(
    layers=layers4,
    max_epochs=50,
    update=adam,
    update_learning_rate=0.0001,
    objective=regularization_objective,
    objective_lambda2=0.009,
    train_split=TrainSplit(eval_size=0.25),
    verbose=3,
)

```

```

In [254]: net1.fit(X, y)

```

```

# Neural Network with 231127 learnable parameters

```

```

## Layer information

```

#	name	size
0	input0	1x2560x1
1	dense1	90
2	dense2	7

epoch	train loss	valid loss	train/val	valid acc	dur
1	3.82806	3.43833	1.11335	0.39198	0.22s
2	3.32719	3.16184	1.05230	0.44757	0.21s
3	2.95896	3.02766	0.97731	0.50922	0.21s
4	2.80678	2.94084	0.95441	0.54715	0.20s
5	2.68662	2.88034	0.93274	0.57194	0.21s
6	2.57092	2.83540	0.90672	0.57845	0.20s
7	2.47766	2.79376	0.88685	0.59068	0.20s
8	2.40327	2.75142	0.87346	0.60676	0.20s
9	2.33777	2.72522	0.85783	0.60676	0.20s
10	2.27657	2.69274	0.84545	0.60981	0.20s
11	2.22343	2.66558	0.83413	0.61332	0.20s
12	2.17331	2.63177	0.82580	0.62074	0.20s
13	2.13015	2.60986	0.81620	0.62294	0.20s
14	2.08821	2.58349	0.80829	0.63472	0.20s
15	2.04784	2.56197	0.79932	0.63342	0.21s
16	2.01216	2.53521	0.79369	0.64474	0.21s
17	1.97773	2.51642	0.78593	0.64213	0.20s
18	1.94520	2.49504	0.77963	0.64519	0.19s
19	1.91394	2.47481	0.77337	0.64825	0.20s
20	1.88388	2.45191	0.76833	0.65345	0.20s
21	1.85673	2.43207	0.76344	0.66087	0.21s
22	1.82998	2.41927	0.75642	0.65911	0.21s
23	1.80263	2.39953	0.75124	0.66393	0.20s
24	1.77815	2.38280	0.74625	0.66002	0.22s
25	1.75281	2.36685	0.74057	0.66568	0.20s
26	1.73016	2.34888	0.73659	0.66653	0.21s
27	1.70754	2.33762	0.73046	0.66698	0.31s
28	1.68457	2.32141	0.72567	0.66478	0.22s
29	1.66362	2.30642	0.72130	0.67259	0.20s
30	1.64336	2.29039	0.71750	0.67389	0.22s
31	1.62443	2.28123	0.71209	0.67259	0.22s
32	1.60414	2.26628	0.70783	0.67259	0.24s
33	1.58505	2.25403	0.70321	0.67389	0.30s
34	1.56690	2.23681	0.70051	0.68130	0.31s
35	1.55103	2.22709	0.69644	0.67825	0.22s
36	1.53289	2.21418	0.69231	0.68261	0.21s
37	1.51545	2.20517	0.68722	0.67695	0.22s
38	1.49841	2.18691	0.68517	0.68215	0.27s
39	1.48423	2.17666	0.68188	0.68436	0.32s
40	1.46786	2.16570	0.67778	0.68436	0.30s
41	1.45190	2.15946	0.67234	0.68612	0.21s
42	1.43553	2.14364	0.66967	0.68917	0.20s
43	1.42174	2.13122	0.66710	0.68872	0.21s
44	1.40805	2.12001	0.66417	0.69178	0.22s
45	1.39390	2.11612	0.65871	0.68436	0.20s
46	1.37806	2.10391	0.65500	0.68697	0.20s
47	1.36520	2.09071	0.65298	0.69087	0.20s
48	1.35196	2.07819	0.65055	0.69132	0.20s
49	1.33972	2.07609	0.64531	0.68957	0.19s
50	1.32508	2.06473	0.64177	0.68827	0.21s


```

f.write('Id,Prediction\n')
index = 1
for pred in cls_res_list:
    f.write('%d,%d\n'%(index, pred))
    index += 1
while index<=1253:
    f.write('%d,0\n'%(index))
    index+=1

```


[7, 7, 4, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 4, 7, 1, 4, 7, 4, 7, 7, 7, 7, 4, 7, 4, 7, 4, 4, 4, 2, 5]

```

In [258]: from nolearn.lasagne.visualize import plot_loss
          from nolearn.lasagne.visualize import plot_conv_weights
          from nolearn.lasagne.visualize import plot_conv_activity
          from nolearn.lasagne.visualize import plot_occlusion
          plot_loss(net1)
          #plot_conv_activity(net1.layers_[1], X[0:1])

```

Out[258]: <module 'matplotlib.pyplot' from '/Users/zexuanwang/anaconda/lib/python2.7/site-packages/matplotlib/pyplot.py'>



```

In [ ]:
In [ ]:
In [ ]:
In [ ]:

```